Plan B: The Show Must Go On

Noctis Zhang

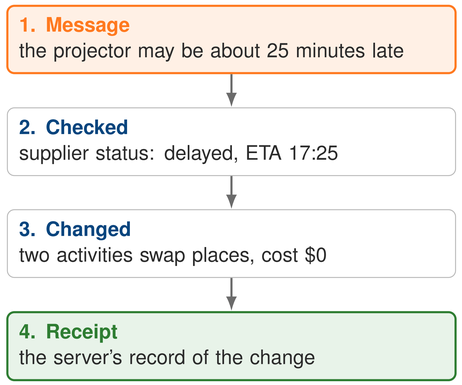
University of Massachusetts Boston • PEAAII Hackathon, October 2026

Plan B decides what to verify, what to prepare, and when to commit a change

Tested on 12 held-out cases (120 runs): no executed violation, no change on the 3 cases that needed none, and no clear winner.

A message is a claim, not a fact

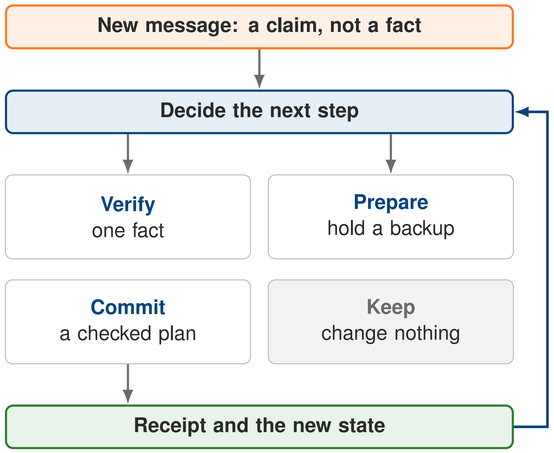
Twenty minutes before a campus showcase opens, a message says the projector may be late.



Practice case dev01, rule baseline, automated run.

In another practice case, a stale reminder and its forwarded copy **change nothing**.

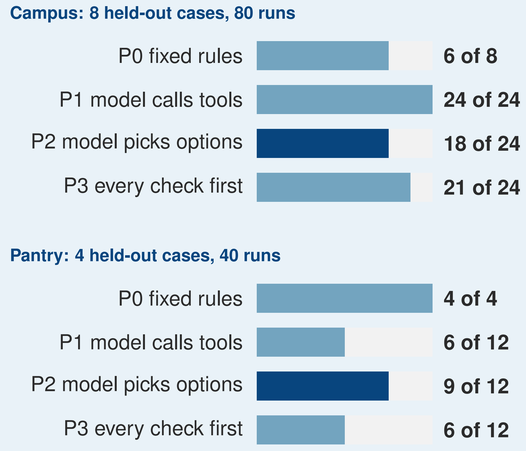
Model chooses, server checks



A language model picks one option. The server built the options, checks every action, and writes the **receipt**. Spending needs a person.

Tested on cases it never saw

Frozen first, then **held-out cases**. Runs that met their outcome:



No consistent advantage for my P2. P1 met every campus run, at three times P2's model cost.

Same loop, a new rule

In a fictional pantry, late stock does not count, and only staff record a handoff or a distribution. **No kit was counted twice** in any held-out run.

What I built

A local server and board with a checked executor, a planner, holds, receipts, four policies and a frozen test harness.

Limits

Fictional sandbox, 12 held-out cases, one model, one reported run after a voided first attempt. My P2 missed cases through a frozen rule, choices that used up the 8 model calls, and a planner defect. No community pilot yet.

Fictional local sandbox. No real bookings, stock or partners.

github.com/nocdoggo/COLLIE (branch hackathon/peeaai-plan, apps/plan_b)