Plan B: The Show Must Go On

Noctis Zhang

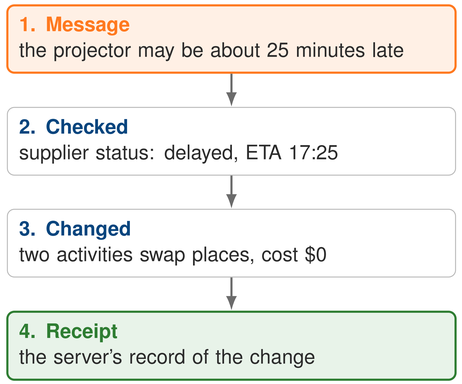
University of Massachusetts Boston • PEAAII Hackathon, October 2026

An AI event coordinator that decides what to verify, what to prepare, and when to commit a change

Version 1.1 on 12 held-out cases (120 runs): no executed action broke a limit, no clear winner. Version 1.2 on 6 fresh held-out cases: 59 of 60 runs met.

A message is a claim, not a fact

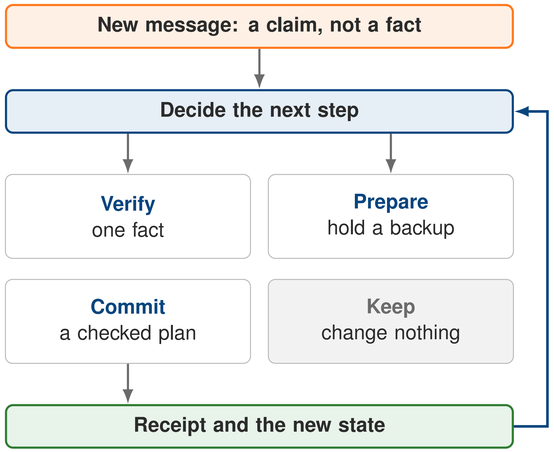
Twenty minutes before a campus showcase opens, a message says the projector may be late.



Practice case dev01, an automated P0 run.

In another practice case, a stale reminder and its forwarded copy **change nothing**.

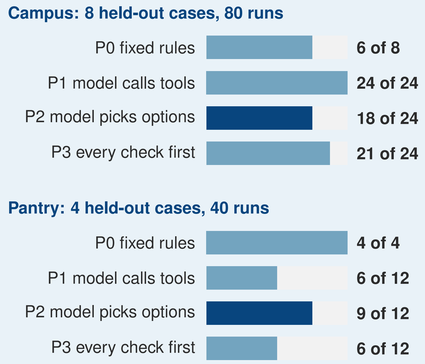
Model chooses, server checks



A language model picks one option. The server builds the options, checks every action and writes the **receipt**. Spending needs a person.

Tested on cases it never saw

Version 1.1 frozen, then **held-out cases**. Runs met, my P2 in dark:



No consistent advantage for my P2. P1 met all 8 campus cases, at three times P2's model cost. **Version 1.2** on 6 fresh cases: P1 to P3 met every run, too few to rank.

Same loop, a fictional pantry

In a fictional pantry, late stock does not count and only staff record a handoff. **No kit was counted twice** in any held-out run.

What I built

A local server and board, a checked executor, a planner, four policies and a frozen test harness.

Limits

Fictional sandbox, 12 held-out cases and 6 fresh ones for 1.2, one model, a voided first attempt. In 1.1 my P2 missed cases through a frozen rule, its 8-call limit and a planner defect. A 1.2 rerun on spent cases is diagnostic only. No pilot yet.

Fictional local sandbox. No real bookings, stock or partners.

github.com/nocdoggo/COLLIE (branch hackathon/peeaai-plan, apps/plan_b)