UMass Boston

# Plan B, The Show Must Go On



**Check the fact**



**Pick only what fits**



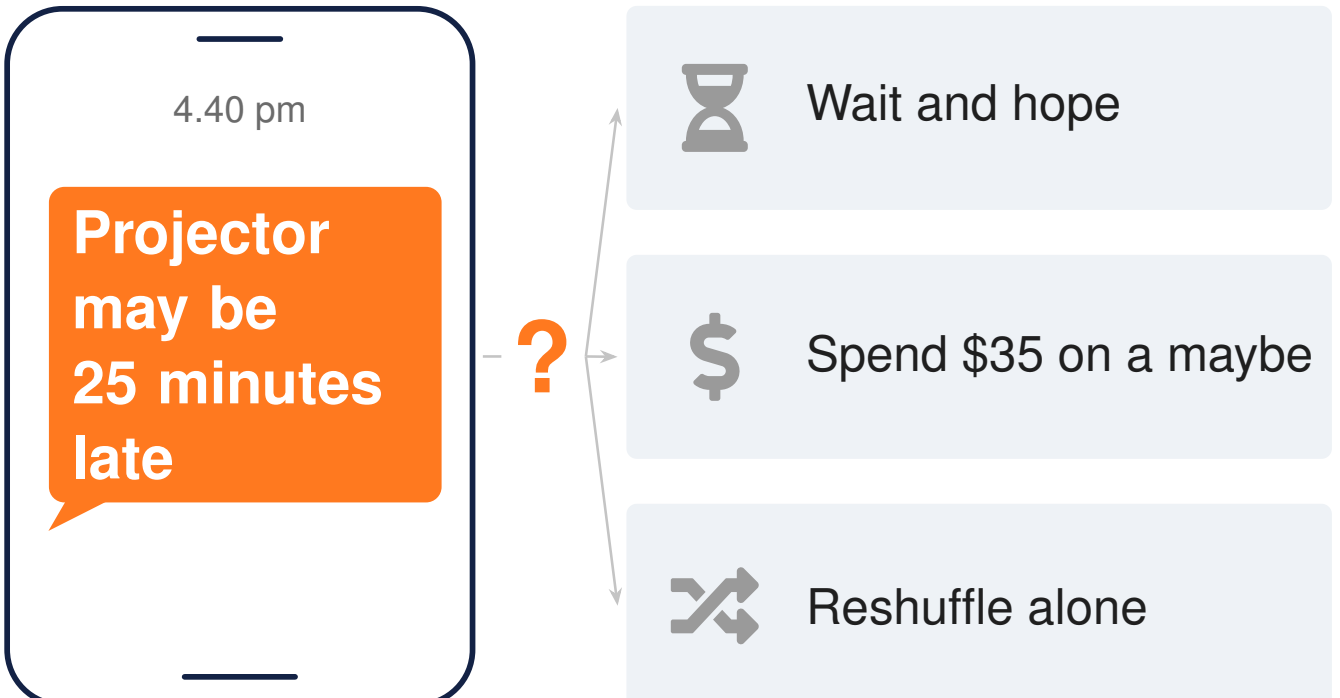
**Leave a receipt**

**Noctis Zhang and Zhen Lu**

Department of Computer Science, University of Massachusetts Boston

PEAAII AI Applications Hackathon 2026, October 2026

## Twenty minutes before doors open, Jordan gets a text




Jordan, a volunteer organizer we made up.

## Plan B rests on three simple ideas

1



### Check the fact

1 look-up, 2 minutes

2



### Pick only what fits

16 changes, 3 fit

3



### Leave a receipt

4.43 pm, $0 spent

**Money waits for a person's OK**

## First, Plan B checks the one fact that matters

**25 minutes late?**



looks up the delivery

**Delayed, about 5.25 pm**

False alarm

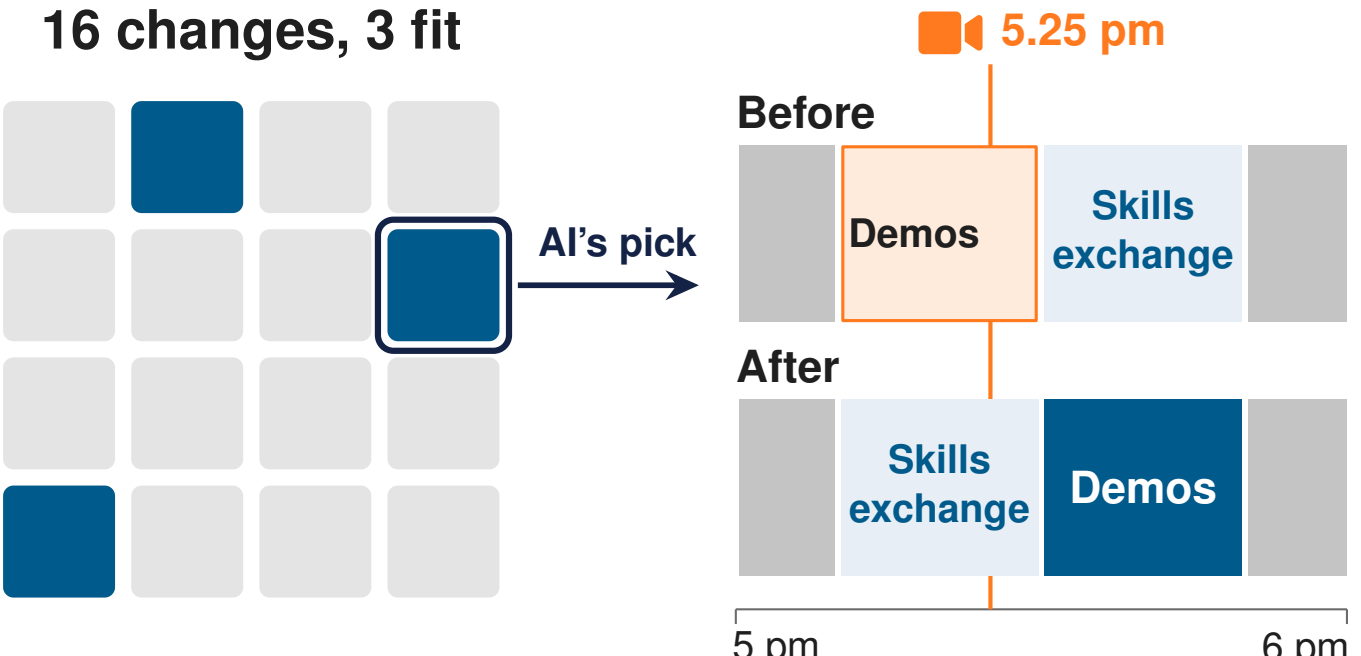
Other order

Repeat

→ **Plan kept**

A message is a claim, not a fact.

# THEN THE AI PICKS ONLY FROM CHANGES THAT FIT




The AI chooses. Plan B makes the change.

# Every change leaves a receipt, and money waits for a person's OK

## Receipt

**4.43 pm**

Swap demos and skills exchange

Version 1 to 2

**$0 spent, $35 not booked**



**Doors open on time**

# THIS IS THE REAL BOARD AFTER THE SWAP

**Before and after**

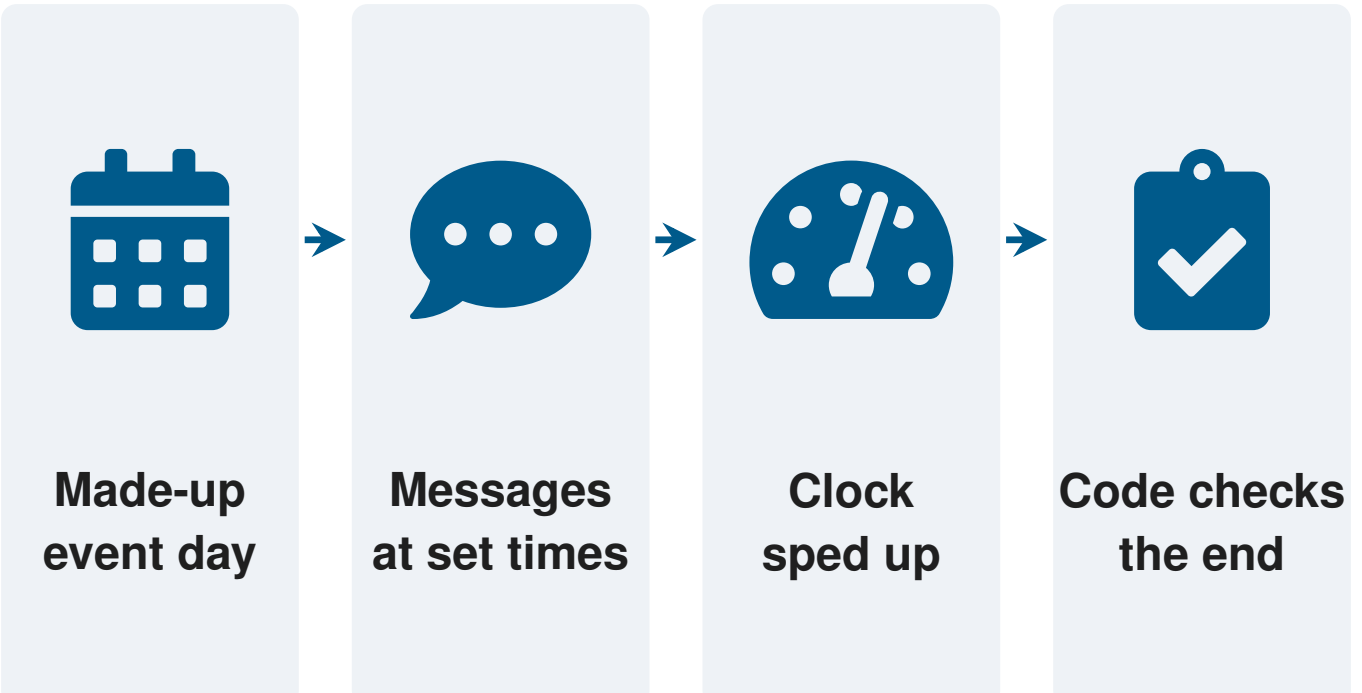
| ACTIVITY | BEFORE (V1) | AFTER (V2) |
|---|---|---|
| **Prepare the event** | 16:30–16:50<br>Fixed community room | 16:30–16:50<br>Fixed community room |
| **Welcome** | 17:00–17:10<br>Fixed community room, Event host | 17:00–17:10<br>Fixed community room, Event host |
| **Community demonstrations** | 17:10–17:30<br>Fixed community room, Primary projector delivery, Demonstration lead | 17:30–17:50<br>Fixed community room, Primary projector delivery, Demonstration lead |
| **Skills exchange** | 17:30–17:50<br>Fixed community room, Skills facilitator | 17:10–17:30<br>Fixed community room, Skills facilitator |
| **Closing** | 17:50–18:00<br>Fixed community room, Event host | 17:50–18:00<br>Fixed community room, Event host |

**Everything else unchanged**

**Two activities swapped**

The running board, in a rehearsal with the fixed rules and no AI.

## We tested Plan B on event days it had never seen



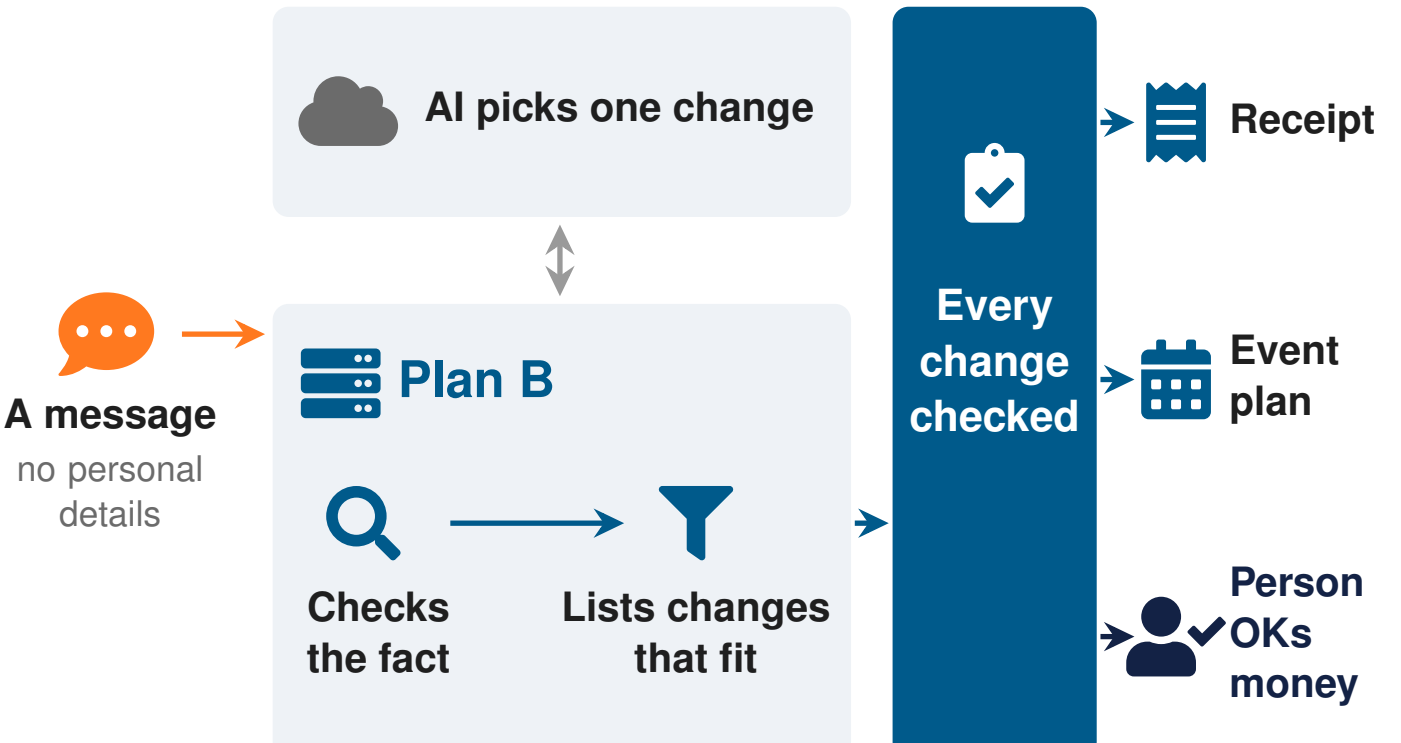

Nothing real is booked or bought.

## 59 OF 60 RUNS MET THE EXPECTED OUTCOME




Every AI run met it. No executed action broke a rule.

# Inside Plan B, careful with messages and money



## Built for student clubs and volunteer groups



### Club officers

100+ UMass Boston groups[28]



### Volunteer coordinators

A helper runs late



### Pantry staff

20 of 20 runs met

made-up pantry days

It checks the text before you spend club money.

# Other tools help, and Plan B works beside them

## Indirect, today

- Group chat[11]
- Shared sheet[12]
- Sign-up sheet[17]
- Campus portal[15]

## Direct, change plans

- Sessionboard[1]
- Cvent[2]
- Bizzabo[3]
- AI assistants[4,5,6]

**Plan B**

From their own pages.

# Plan B is safer, more flexible and more personal



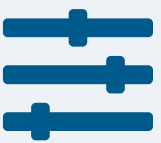
## Safer

The three ideas, built in



## More flexible

3 kinds of event, your AI or none



## More personal

Your event, your spending limit

We have not tested the others.

# Check the fact, pick only what fits, leave a receipt

1



**One look-up**

2



**The free swap**

3



**Doors on time, $0**

**Money waits for a person’s OK**

**Code**
github.com/nocdoggo/FA26-COLLIE

Next, a walk-through with staff.

**Thanks, PEAAII organizers**

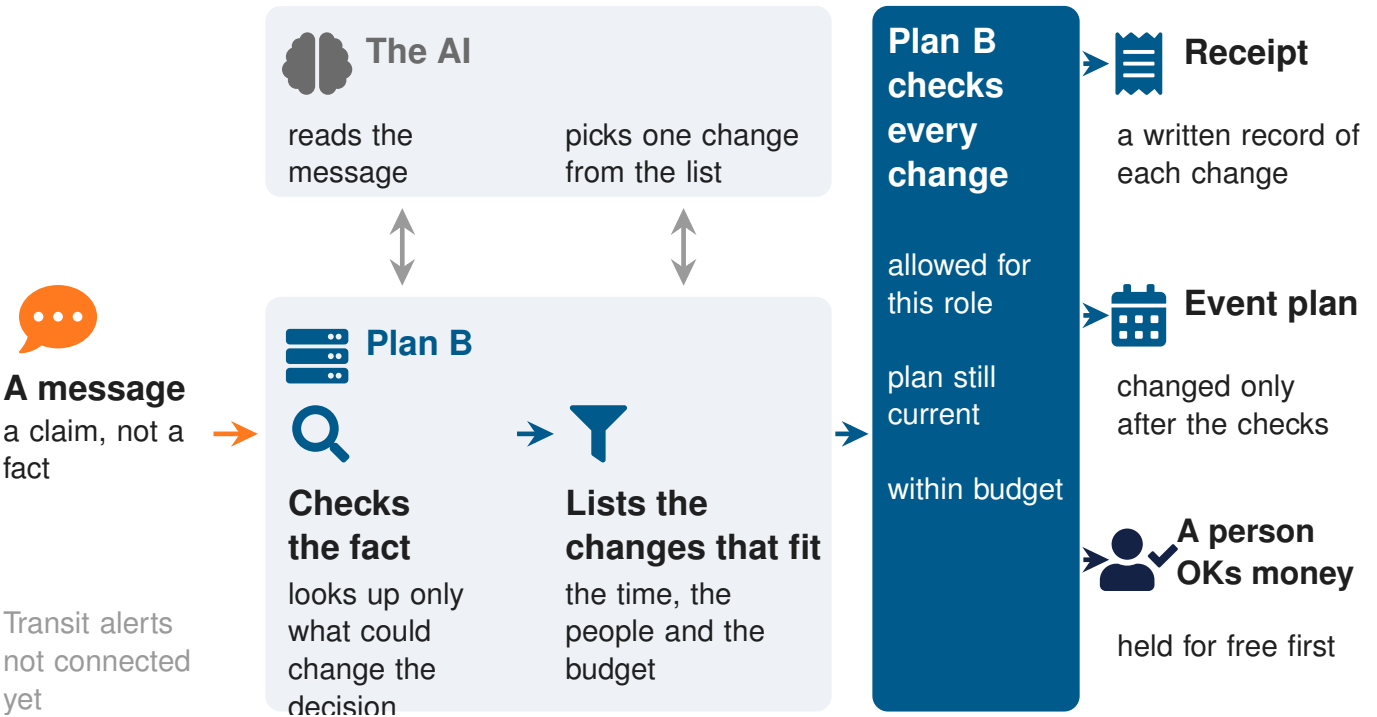
# More detail, the six sealed test cases

| Case | What happens | Met means | No AI | With AI |
|---|---|---|---|---|
| **Projector late** | The projector is really late, past its slot, so the swap does not work | a ten minute shift, or the spare after a person's OK | 1 of 1 | 9 of 9 |
| **False alarm** | An anonymous call says the projector will be late. The delivery is on time | no change at all | 1 of 1 | 9 of 9 |
| **Volunteer late** | The skills exchange leader texts that their bus is out of service | the swap, so every activity still happens | **0 of 1** | 9 of 9 |
| **Cannot be saved** | The projector fails its test, and the spare cannot be confirmed in time | Plan B says plainly it cannot be fully done | 1 of 1 | 9 of 9 |
| **Pantry delivery** | A late delivery cannot arrive in time, and a partner's kits can | 40 kits handed out by the deadline | 1 of 1 | 9 of 9 |
| **Pantry claims** | Wrong and stale delivery claims, and a partner who declines | only kits really on the shelf counted, the shortfall said plainly | 1 of 1 | 9 of 9 |

Each case ran once with the fixed rules and three times with each of three AI approaches (our coordinator, the check-all variant, the tool-calling model). The fixed rules were never written to read the volunteer's own message. In that case no approach looked anything up, all took the volunteer's word. The six cases use new wording but the same kinds of trouble as earlier cases.

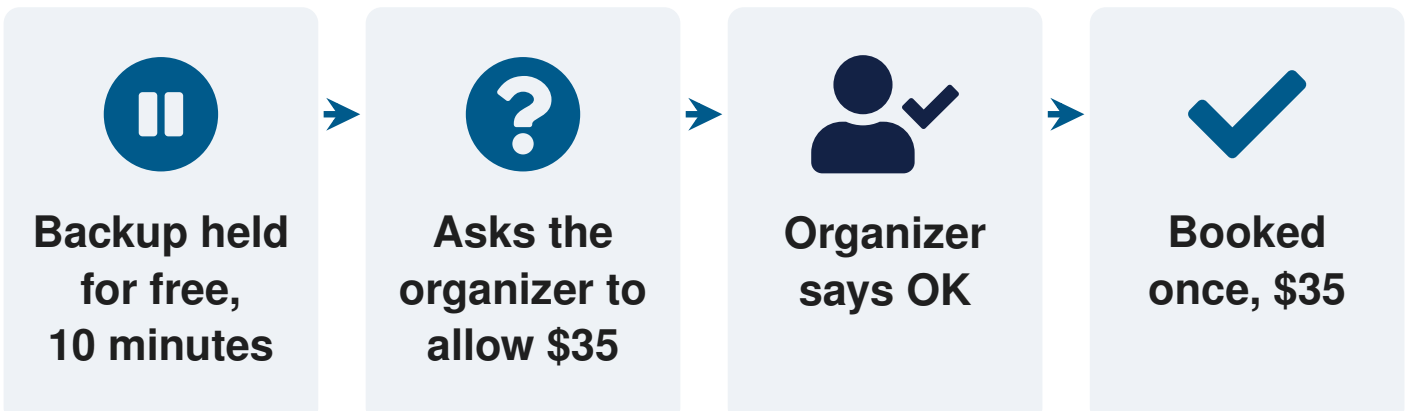
Source, the final test's results folder in the code repository, with its notes on the one miss.

# More detail, how Plan B works inside



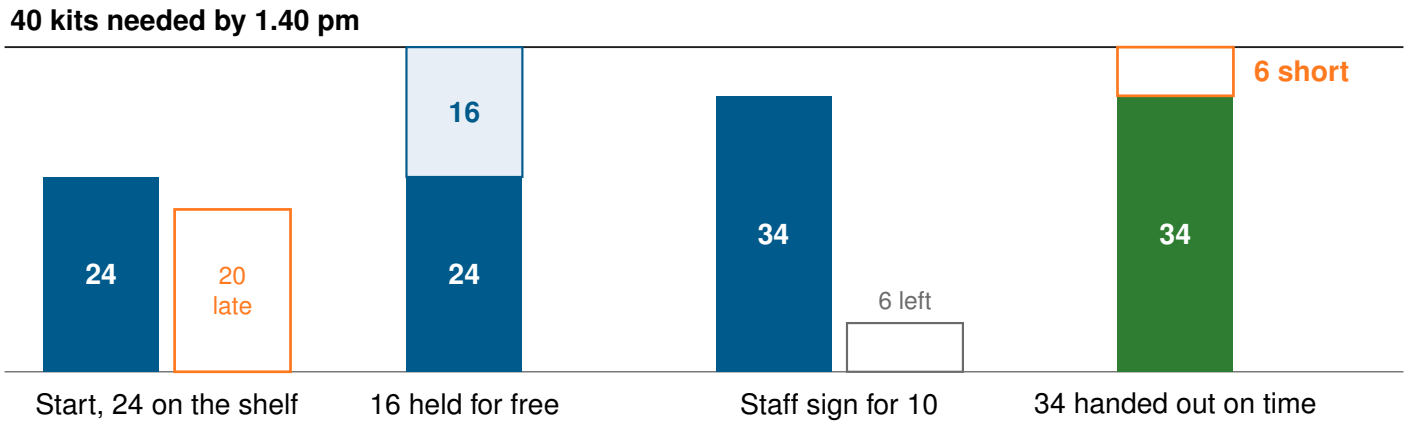

Only Plan B's server can make a change, in every approach tested. In our coordinator the AI picks from the changes that fit.

# What if the projector came at 5.45 pm




If the look-up had said 5.45 pm, too late for either slot, Plan B holds the backup projector for free, kept aside for ten minutes and released on its own. Its spending allowance is $0, so it asks the organizer to allow $35. Only after the OK is the backup booked, once. In a recorded run of this case the AI made that choice at 4.48 pm. This is a separate practice case, not part of Jordan's evening.

# THE SAME IDEAS AT A MADE-UP FOOD PANTRY




40 kits are needed by 1.40 pm. 24 are on the shelf and 20 are late, so the late ones are not counted until staff sign for them.

Plan B holds a partner's 16 kits for free and asks staff for the handoff. Staff sign for 10.

34 kits are handed out on time. Plan B tells staff 6 will be short, and staff decide what to tell the line. No kit is counted twice.

**Plan B counts kits, never people. Who gets food is always a person's call.**

Made up, not based on any real pantry's data. Identical boxed kits only.

# The earlier 12 case test, and what it found

| Runs that met | Campus | Pantry |
|---|---|---|
| The fixed rules | 6 of 8 | 4 of 4 |
| The tool-calling model | 24 of 24 | 6 of 12 |
| Our coordinator | 18 of 24 | 9 of 12 |
| The check-all variant | 21 of 24 | 6 of 12 |

Twelve sealed cases, 120 runs, version 1.1 of the code. Source, the earlier test's notes in the code repository.

The earlier test, on twelve sealed cases, found four problems. A partner's stock that was already checked as unavailable stayed an option, a message about a person running late was not treated as important, the AI was asked again about choices it had already passed on, and a choice refused as out of date cost a new AI call.

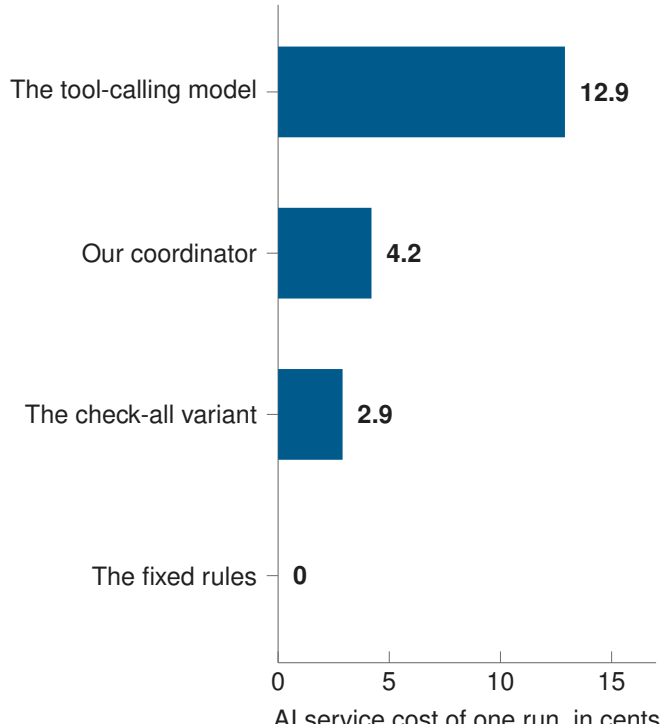
I worked on all four before the final test. A later rerun of the earlier cases still misses two of the cases these changes were aimed at, so not every problem is fully solved.

# WHO ELSE DOES THIS, FROM THEIR OWN PAGES

| Tool | What its own page says it is for | What Plan B adds for a small group |
|---|---|---|
| **Sessionboard [1]** | AI agenda building, checks agenda changes against rules and sends big changes for sign-off | the one fact checked first, a receipt for every change, free and open |
| **Cvent [2], Bizzabo [3]** | agents for event teams, announced in 2026, whose changes a person reviews or approves | runs today on one laptop, free, with no vendor contract |
| **ChatGPT agent, Copilot, Gemini [4, 5, 6]** | general assistants that ask before purchases and other actions with real-world consequences | picks only from changes that fit this event's rules |
| **Zapier Agents [7]** | automations that stop for approval when the user adds that step | the checks come built in |
| **Group chat, shared sheet, sign-up sheet, campus portal [11, 12, 15, 17]** | where the text lands today | works beside them, not instead of them |

We have not tested the others. Full sources in the project documentation.

# What the final test cost, and the limits per run




All 60 runs cost **$3.61** in AI service fees, with 194 AI calls and no service errors.

Limits per run, at most 48 tool calls, 8 AI calls, $0.50 of AI service fees and 900 seconds. By default the AI service's bill is capped at $20 in total. The final test ran under a $40 cap.

These limits cap the AI service's bill, not event money. Event money always waits for a person's OK.

**The same outcomes cost about a third as much when the AI picks from the changes that fit.**

Source, the final test's results and spend notes, and the default settings in Plan B's code.

# THANKS, AND THE SOURCES CITED IN THIS DECK


## Thanks

To the PEAAII organizers at UMass Boston. The AI model is Claude Sonnet 5.5 by Anthropic, reached through OpenRouter. The deck uses the elegant beamer template by Lars Spreng, and the poster the Gemini theme by Anish Athalye with the UMass Boston colors. Plan B is built with FastAPI, uvicorn and Pydantic. Thanks do not imply that any of them endorses Plan B.


## Code

github.com/nocdoggo/FA26-COLLIE, free and open source (GPL-3).

## Sources


[1] Sessionboard, AI Agenda Builder and Event Scheduler, 2026.
[2] Cvent press release, July 14, 2026.
[3] Bizzabo, Agentic Operating System for Events, October 12, 2026.
[4] OpenAI, Introducing ChatGPT agent.
[5] Microsoft Support, Browse with Copilot.
[6] Google, Manage your tasks and workflows in Gemini Apps.
[7] Zapier Help Center, Add approval steps to your agent's instructions.
[11] Microsoft Support, What is GroupMe.
[12] Google Workspace, Sheets.
[15] UMass Boston, Student Activities.
[17] SignUpGenius, Online Sign Up Sheets.
[28] UMass Boston, Student Groups and Activities, checked October 26, 2026.
[33] OWASP, LLM01 2025 Prompt Injection.
[37] NCSC, Prompt injection is not SQL injection, December 8, 2025.
Full list in the project documentation.